

2 Summarizing networks using statistics

Regardless of which representation we use, networks are sufficiently complicated objects that we almost always need to use statistical methods in order to gain an intuitive sense of their shape and variations, or to describe their patterns, or test whether these patterns are meaningful. In this section, we will survey some of the most basic of these descriptive statistics. These statistics fall largely into three categories:

1. **connectivity** measures are directly or indirectly related to the number of connections a node has, and to the overall degree structure of the network,
2. **motif** measures count the frequency of specific subgraphs in a network,
3. **positional** measures are related to where in the network a node sits, and to the distances between nodes in the network.

Many of these network measures come in two flavors:

- **node-level** or “local” measures, which we calculate for an individual node i , and
- **network-level** or “global” measures, which we calculate across the entire network.

Sometimes, we can obtain the corresponding network-level statistic of a particular quantity by simply computing the average node-level quantity, but not always.

Because network representations can vary, so too do the algorithms for calculating specific network statistics. As a result, before beginning any network analysis, we need to first identify the graph properties of representation being used: Is the network simple? Are edges directed? Are edges weighted? Do nodes have metadata? Does the graph have one or several components? Etc.

2.1 Node degree

The most basic of all network statistics is the **degree** of a node, which counts the number of connections it has in the network. The degree structure of a network is the *first-order* description of a network, and is often sufficient to drive many other statistical patterns in the network's structure. Later, we'll learn how to assess the extent to which degrees alone can explain other structural patterns in networks.

2.1.1 Simple networks

When edges are undirected, an edge (i, j) contributes to the degree of both i and j . By convention, k_i denotes the degree of vertex i . For a simple adjacency matrix A , the degree of node i is

$$k_i = \sum_{j=1}^n A_{ij} = \sum_{j=1}^n A_{ji} . \quad (1)$$

That is, the degree k_i is just the sum of the i th column, or i th row, of the adjacency matrix A . The row- and column-sum equivalence holds only because edges are undirected. If A represents a weighted network, this sum is called the node **strength**; the term “degree” is reserved for unweighted counts.⁵

Every edge in an undirected network contributes twice to some degree (once for each endpoint or “stub”), and so the sum of all degrees in a network must be equal to twice the total number of edges in a network m :

$$m = \frac{1}{2} \sum_{i=1}^n k_i = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n A_{ij} = \sum_{i=1}^n \sum_{j>i}^n A_{ij} . \quad (2)$$

Using Eq. (2), we can derive a simple expression for the network's mean degree $\langle k \rangle$:

$$\langle k \rangle = \frac{1}{n} \sum_{i=1}^n k_i = \frac{2m}{n} . \quad (3)$$

And finally, if we divide the number of edges m by maximum number of possible edges in a simple network, or divide the mean degree by its largest possible value, we obtain a quantity that is sometimes called the network's **connectance**⁶ or **density**:


$$\rho = \frac{m}{\binom{n}{2}} = \frac{2m}{n(n-1)} = \frac{\langle k \rangle}{n-1} . \quad (4)$$

⁵If A is a non-binary adjacency matrix, then a simple row sum will yield the strength s_i of node i ; if the summation is modified as a count of non-zero entries, we will get back the correct value of k_i .

⁶Sometimes, connectance is defined as $\langle k \rangle / n$, which is asymptotically equivalent to $\langle k \rangle / (n-1)$.

2.1.2 Directed networks

In a directed network, a node i has three types of degrees: its **in-degree** k_i^{in} , meaning the number of connections to i ; its **out-degree** k_i^{out} , the number of connections from i ; and its **total degree** k_i^{tot} , the number of neighbors, ignoring edge directionality. The latter is equivalent to treating directed edges as undirected, and the “collapsing” multi-edges so that only one connection (i, j) exists.



On a binary adjacency matrix A representing a directed network, the in-, out-, and total degrees are defined as

$$k_i^{\text{out}} = \sum_{j=1}^n A_{ij} \quad (5)$$

$$k_i^{\text{in}} = \sum_{j=1}^n A_{ji} \quad (6)$$

$$k_i^{\text{tot}} = \sum_{j=1}^n A_{ji} \vee A_{ij} , \quad (7)$$

where \vee is an *or* function. In a directed network, the number of edges m is the sum of the entire adjacency matrix, since each directed edge appears exactly once in A . The fact that each edge contributes exactly once to some in-degree and once to some out-degree also implies that the average in-degree equals the average out-degree:

$$\langle k^{\text{out}} \rangle = \frac{1}{n} \sum_{i=1}^n k_i^{\text{out}} = \frac{1}{n} \sum_{i=1}^n \left(\sum_{j=1}^n A_{ij} \right) = \frac{m}{n} = \sum_{i=1}^n \left(\sum_{j=1}^n A_{ji} \right) = \frac{1}{n} \sum_{i=1}^n k_i^{\text{in}} = \langle k^{\text{in}} \rangle . \quad (8)$$

2.2 Degree distributions


If we calculate the degrees of all the vertices in a network, and then write them out, we get something called a **degree sequence**, which enumerates the network’s degree structure. For the simple graph in Figure 1, its degree sequence is

$$(1, 2, 2, 2, 3, 4) .$$

But a simple sequence doesn't provide much information about the *variability* of node degree across the network. To capture this behavior, we use something called the **degree distribution**, denoted $\Pr(k)$, which is simply the probability that a vertex selected uniformly at random will have k neighbors. A degree distribution is a normalized histogram of the degree values. For the simple graph in Figure 1, its degree distribution is

k	1	2	3	4
$\Pr(k)$	1/6	3/6	1/6	1/6

where $\Pr(k) = 0$ for all other values of k . The standard way to present a degree distribution $\Pr(k)$ is via a bar plot, like so:



2.2.1 Exploring the degree distribution


Most real-world networks, biological or otherwise, exhibit a “heavy-tailed” degree distribution, meaning that the variance⁷ of the degree distribution σ^2 is large relative to its mean $\langle k \rangle$. Because degrees must be non-negative (they are “count” variables), this behavior implies that the degree distribution $\Pr(k)$ is “right-skewed,” meaning that the largest degree vertices are often several orders of magnitude better connected than the typical node.

This kind of high variance, however, can be difficult to see in a simple distribution plot, and instead, it is conventional to (i) use a log-log plot, which visually compresses the variance in both k and $\Pr(k)$, and (ii) plot the *complementary cumulative distribution function* or CCDF, defined as $\Pr(K \geq k)$. In this expression, K is now a random variable, and the CCDF plots the fraction of nodes with degree *at least* some value k . If the minimum degree of a network is $k = 1$, then $\Pr(K \geq 1) = 1$, because all nodes have degree at least the minimum.⁸


⁷Recall that the sample variance of the degrees would be defined as $\sigma^2 = \frac{1}{n-1} \sum_{i=1}^n (k_i - \langle k \rangle)^2$.

⁸Applying a `cumsum` function to the PDF of a degree distribution *almost* yields the CDF, which is defined as $\Pr(K < k)$. Do you see why? Think about what happens to the first element under `cumsum` and whether that is correct for $\Pr(K < \min(k_i))$. After correcting for this behavior, one can obtain the CCDF by noting that $\Pr(K \geq k) = 1 - \Pr(K < k)$.

To illustrate how the log-plots change the kind of information we can see in a degree distribution, and how plotting the CCDF improves the visual interpretability, we'll use a concrete example. Below is a network of $n = 1490$ weblogs from around the 2004 U.S. Presidential election, along with a standard plot of the degree distribution, which spans $m = 19090$ edges and has a mean degree $\langle k \rangle = 25.6$.⁹ The details of how these data were collected, or what they mean is not that important, except to note that this is a web graph, and is thus a directed network. Notice that the degree distribution emphasizes the low-degree portion of the distribution, for $k \leq 10$ or so. In fact, the largest degree is $k_{\max} = 351$, a value more than 13 times larger than the mean, but you can't tell that from the plot.




What follows is a sequence of plots, showing the same degree data, as (i) $\log_{10} \Pr(k)$ vs. k (a semilog-y plot), then (ii) $\log_{10} \Pr(k)$ vs. $\log_{10} k$ (a loglog plot), and finally (iii) $\log_{10} \Pr(K \geq k)$ vs. $\log_{10} k$ (a loglog plot of the CCDF).




In the semilog-y plot above, we can now see more of the low frequency / high degree events, which we call the “upper tail” of the distribution. The dots represent nodes with unique degree values; hence, $\Pr(k) = 1/n$ for every unique value of k , which is the lowest value possible in $\Pr(k)$. This

⁹The visualization comes from Karrer & Newman, *Phys. Rev. E* **83**, 016107 (2011), and the network comes from Adamic & Glance, *Proc. WWW-2005 Workshop on the Weblogging Ecosystem* (2005).

minimum value is what causes the “flat-bottom” structure of the degree distribution, which starts just over $k = 50$.




The loglog plot above stretches out the low-degree structure while compressing the range of the high-degree portion of the distribution, allowing us to see both parts together. But, the $1/n$ artifact is still apparent, and the jagged structure at $k > 60$ or so is too noisy to see any clear pattern. The solution is the CCDF, which is a monotonically decreasing function that smooths out many of the jaggies in the upper tail.



In the above plot, we can now see clearly that the distribution tends to bend down around $k = 64$, falling off more quickly than it did below that point. Moreover, we can read off, directly from the plot, the fact that 90% of nodes have degree $k \leq 67$ (at what values of x does a horizontal line at $y = 10^{-1}$ intersect the CCDF?), and only 1% of nodes have degree $k > 169$ (at what value of x does a horizontal line at $y = 10^{-2}$ intersect the CCDF?). Reflecting this extremely uneven distribution of edges, the bottom 90% of nodes, by degree, touch only 53% of all edges, while those top 1% touch 10% of all edges. This dramatic inequality is ubiquitous in real network degree distributions, and highlights the important role that these “high degree” nodes play in structuring the network. The clarity provided by the CCDF on a loglog plot makes it the most useful way to visualize degree distributions.

How much variability do degree distributions exhibit in reality? The figure below shows the results of applying the above CCDF analysis to 100 networks drawn from biological, social, information, and technological networks. The answer is: they vary enormously in their shape, but are generally high variance, even “heavy tailed.” In other words, real-world networks structurally diverse, even in something as simple as their degree structure, and the only commonality across degree distributions is their often enormously high variance.



2.3 Small motifs

A **motif** is a small, connected subgraph $G_o = (V_o \subset V, E_o \subset E)$ with a particular pattern of edges among a small set of nodes, and usually, we want to count how often a motif G_o occurs within a larger network G . The computational cost of counting motifs scales up quickly with the size of V_o , and thus most motif-based measures focus on the frequencies of small motifs, e.g., 2, 3, or 4 nodes¹⁰

2.3.1 Reciprocity (directed networks) Skip - Not covered

The smallest non-trivial motif in a directed network is a reciprocated edge, i.e., a *2-cycle* or a cycle of length exactly two, and we use a measure called **reciprocity** r to quantify their prevalence.

¹⁰Computational complexity and *subgraph isomorphisms* are key problems for motif-counting algorithms. Suppose we are searching for a motif on n_o nodes. If we blindly check all $f(n_o) = n^{n_o}$ groups of nodes, each found motif will occur $|n_o|!$ times in this search, due to isomorphisms. That is, there are $|n_o|!$ ways to permute the node labels of any unique motif. Asymptotically, how expensive is such a search for $n_o = \{2, 3, 4, 5, 6\}$?

Reciprocity comes in both node-level (local) and network-level (global) versions.



The prevalence of this motif is often interpreted functionally. In some ecological networks, reciprocity can indicate the degree of mutualism present in an ecosystem, while in gene regulatory networks, reciprocated links indicate the presence of self-regulating interactions.

The reciprocity of a network r is given by the fraction of links that are reciprocated, meaning it has a maximum value of $r = 1$. To compute r , we simply count the number of 2-cycles and divide by the number of edges. When edges have unit weight, a network's reciprocity is defined as

$$r = \frac{1}{m} \sum_{ij} A_{ij} A_{ji} . \quad (9)$$

Note that in a directed network, m counts the number of directed edges, and hence we normalize by $m = \sum_{ij} A_{ij}$ to place r on the unit interval $[0, 1]$.


We can also quantify the reciprocity that an individual vertex experiences, which is a node-level measure called *local reciprocity*. Instead of summing over the entire network, in this case, we count the number of 2-cycles attached to vertex i , and divide by its out-degree:

$$r_i = \frac{1}{k_i^{\text{out}}} \sum_j A_{ij} A_{ji} . \quad (10)$$

2.3.2 Clustering coefficient (undirected networks)

In an undirected network, the smallest interesting motif is a triangle, and the **clustering coefficient** C is a network-level measure that captures the relative frequency or the density of triangles in a network. Concretely, C quantifies the fraction of “connected triples,” that is, a pair of edges $(i, j), (j, k)$ that are “closed” by an edge (k, i) to form a triangle.

Like reciprocity, C is defined mathematically as a count of the desired motif (a triangle), normalized




by a count of the number of possible motifs:

$$C = \frac{(\text{number of triangles}) \times 3}{(\text{number of connected triples})} \quad (11)$$

$$= \sum_{i=1}^n \sum_{j=1}^n \sum_{k=1}^n A_{ij} A_{jk} A_{ki} \bigg/ \sum_{i=1}^n \sum_{j=1}^n \sum_{k \neq i}^n A_{ij} A_{jk} \quad , \quad (12)$$

where the factor of 3 in Eq. (11) comes from the symmetry of the triangle (every triangle contains three connected triples), and the $k \neq i$ restriction in the denominator of in Eq. (12) prevents counting $(i, j), (j, i)$ as a triple.¹¹ Because the numerator can never exceed the denominator, and they are both count variables, $C \in [0, 1]$, and we say that it measures the density of triangles in the network.¹² Calculating Eq. (11) directly would take $O(n^3)$ time. How long would it take to calculate C using an adjacency list representation?

A worked example. Consider a simple network composed of cycle and a clique:



The clique has 6 vertices, making $\binom{6}{3} = 20$ triangles, and thus also $20 \times 3 = 60$ connected triples. The cycle contains 6 connected triples and no triangles. Finally, there are 2 connected triples starting from the cycle and 5 connected triples starting from the clique that use the edge joining the cycle to the clique. Adding this all up yields a clustering coefficient of $C = 60/73 = 0.82$.

¹¹In linear algebra terms, Eq. (12) is equivalent to $C = \text{trace}(A^3) / \sum_{i \neq j} A^2$. This is because the diagonal of the matrix A^q counts the number of cycles of length q that “originate” at each node.

¹²When the network is fully connected, in which every edge exists, $C = 1$ (do you see why?). Similarly, in any network without triangles, e.g., a tree, a bipartite network, etc., $C = 0$ (do you see why?).

Local clustering coefficient. The clustering coefficient can also be defined as a node-level measure, which captures the fraction of neighbors j, k of node i that are themselves connected by an edge (j, k) :

$$C_i = \frac{(\text{number of pairs of neighbors of } i \text{ that are connected})}{(\text{number of pairs of neighbors of } i)} \quad (13)$$

$$= \sum_{j < k} A_{ij} A_{jk} A_{ki} / \binom{k_i}{2} . \quad (14)$$

In the cycle-plus-clique example above, the highlighted vertex has a local clustering coefficient of $C_i = 0$ because it participates in no triangles. Its immediate neighbor in the clique is a more interesting case, with $C_i = 10/15 = 0.67$.

2.3.3 Feed-forward and feedback loops (directed networks) Skip - Not covered

In directed networks, and particularly in biological networks, two slightly larger motifs of special interest are the feed-forward loop and the feedback loop. These motifs are interesting because they represent basic biological circuitry for governing the flow of activation across molecular networks¹³



A feed-forward loop is defined as one that has a two-step path $(i, j), (j, k)$ that is shortcut by an edge (i, k) that “feeds forward” the signal passing along the longer path. In this case, we would say that node j is the “feed forward” node. Similarly, a feedback loop is defined the same, except that the third edge is (k, i) , which “feeds back” the signal that passed along the longer path.

If a network is directed, then we can count the number of FFL and FBL straightforwardly under an adjacency matrix representation:

$$\text{FFL} = \sum_{i=1} \sum_j \sum_{k \neq i} A_{ij} A_{jk} A_{ik} \quad (15)$$

$$\text{FBL} = \sum_{i < j < k} A_{ij} A_{jk} A_{ki} , \quad (16)$$

¹³For instance, see [https://en.wikipedia.org/wiki/Feed_forward_\(control\)](https://en.wikipedia.org/wiki/Feed_forward_(control)).

These adjacency matrix calculations take $O(n^3)$ time because each has three nested loops over n nodes, and hence don't scale well to large networks.


A more efficient approach, taking $O(n + m\langle k \rangle^2)$ time, would use a simple “path enumeration” algorithm on an adjacency list to find all paths $(i, j), (j, k)$ (where $i \neq k$) and, for each, check whether the edge (i, k) exists, indicating a FFL motif, or (k, i) exists, indicating a FBL motif.

2.3.4 Larger motifs

A motif can be defined on any size subgraph $n_o < n$, and sometimes “large” motifs, when $n_o > 3$, are interesting. However, as n_o increases, it becomes increasingly expensive to both enumerate and count such motifs. For instance, while there are 6 distinct motifs on $n_o = 4$ nodes, $n_o = 5$ nodes have 21, $n_o = 6$ nodes have 112, and $n_o = 7$ nodes have 853. Even if we successfully count the frequencies of a large motif, its interpretation also requires special care, because there can be many different ways to “normalize” that count relative to an expectation, as we did in Eq. (9) for reciprocated links and Eq. (11) for triangles. We will revisit the idea of assessing the frequency of larger motifs when we discuss random graph models.

2.4 Geodesic paths and network position

Recall that a *path* in a network is a sequence of distinct vertices $x \rightarrow y \rightarrow \dots \rightarrow z$ such that for each consecutive pair of vertices $i \rightarrow j$, there exists an edge $(i, j) \in E$ in the network, and the *length* of the path is the number of edges it crosses. A *geodesic* or *shortest path* is the shortest of all possible paths between two vertices. These paths serve as the basis for a number of positional measures of network structure.¹⁴



To calculate nearly any positional measure, we must first calculate the geodesic distances ℓ_{ij} between all pairs of nodes i, j . Usually, we accomplish this by applying an off-the-shelf algorithm for

¹⁴Geodesic paths are fun! The classic *Six Degrees of Kevin Bacon* game asks players to enumerate a path on the bipartite network of actors and the movies they appear in from a selected actor to Kevin Bacon. The popular website <https://oracleofbacon.org> returns answers derived from applying a shortest-path algorithm. If one prefers academics, the website <https://www.csauthors.net/distance/> does the same, but for pairs of computer scientists, including yours truly.

the *All Pairs Shortest Paths* (APSP) problem.¹⁵ The output will be a *pairwise distance* matrix ℓ , where the entry ℓ_{ij} gives the length of the geodesic path between i and j . Crucially, if no such path exists (because i and j are in different *components* of the network, e.g., the singleton node vs. any other node in the example above), then $\ell_{ij} = \infty$, and we say that j is not *reachable* from i .

If G is a weighted network, the pairwise distance matrix ℓ will contain the *weight* of the *lightest* path from i to j . In this case, the lightest-weight path may, in fact, have more edges in it than the shortest length path (ignoring edge weights). In an unweighted network, edges have unit weight, and hence ℓ_{ij} is the number of edges in the path.

The path structure of a network also determines whether it is composed of one or multiple *components*, each of which is a set of nodes T such that for any pair of members $i, j \in T$, there exists a path $i \rightarrow \dots \rightarrow j$ or vice versa. If all nodes are pairwise reachable from each other, then a network has a single component. Else, the network has multiple components, and there is always a *largest connected component* (LCC), i.e., the component with the largest number of nodes in it.¹⁶ In an undirected network, the matrix ℓ is always symmetric, because the shortest path from i to j can be no longer than the shortest path from j to i .

2.4.1 Diameter

The first positional measure is the network's **diameter**, which is the length of the longest of all geodesic paths and is meant to evoke the notion of a volume in a metric space. The diameter of a network is a global measure, and is defined as

$$\text{diameter} = \max_{i,j} \ell_{ij} . \quad (17)$$

If the network is not connected, i.e., if there is more than one component, then the diameter will trivially be infinite. In these cases, it is customary to drop all $\ell_{ij} = \infty$ elements from the calculation, and say the diameter is the longest geodesic among all paths that exist.

2.4.2 Eccentricity

While the diameter is a global measure, the **eccentricity** is an equivalent local measure—it measures the length of the longest path originating at some node i , and is defined as

$$\epsilon_i = \max_j \ell_{ij} , \quad (18)$$

where again, it is customary to drop all $\ell_{ij} = \infty$ elements from the calculation.

¹⁵Most commonly, the Floyd-Warshall algorithm. Alternatively, applying Dijkstra's algorithm, for solving the single-source shortest path problem, once for each node will build the pairwise distance matrix ℓ one row at a time. If the network is weighted, there are some technical restrictions on the weight values.

¹⁶The number and size of components can be computed straightforwardly, by analyzing the output of an APSP algorithm, or simply a breadth-first or depth-first search algorithm. Do you see how?

2.4.3 Mean geodesic path length


The **mean geodesic path length** is network-level (global) measure of the average distance between a pair of uniformly random nodes, and is defined as

$$\langle \ell \rangle = \frac{1}{Z} \sum_{ij} \ell_{ij} , \quad (19)$$

where here it is customary to drop both all $\ell_{ij} = \infty$ elements and $\ell_{ij} = 0$ elements (the diagonal of ℓ) from the calculation, and Z counts the number of the remaining distances.

An example. Consider again the directed network of political blogs from Section 2.2.1, which has $n = 1490$ nodes. Applying an APSP algorithm to the network yields a 1490×1490 pairwise distance matrix ℓ . This network is not connected: it has 268 components, the largest of which contains 1222 nodes (82%), and the smallest of which contains 1 node.

Taking the set of non-zero and non-infinite geodesic distances, the following two plots show the normalized histograms of those distance, once as a simple bar plot and the other as a semilog-y bar plot. The mean geodesic distance is $\langle \ell \rangle = 2.74$, which we can see reflected in the simple bar plot, where the vast majority of distances are 2 or 3. The diameter, however, is not particularly visible in this version. The longest geodesic distance is $\ell_{\max} = 8$, and occurs exactly once, i.e., there is exactly one pair of vertices that are separated by a path of length 8, which is one path out of 1,492,064 pairs of connected nodes. This value is a bit easier to see in the semilog-y plot, which also shows more of a smooth fall-off in the distances above the mean, to the maximum at 8.



Small worlds and network diameter.

In mathematical models of networks, the diameter plays a special role and can often be shown to vary in a clean functional way with the size of the network. If the diameter grows very slowly as a function of network size, e.g., $O(\log n)$, a network is said to exhibit the “small world” property.

The basic idea of “small world” networks comes from a seminal study in social networks by the American sociologist Stanley Milgram (1933–1984).¹⁷ Milgram mailed letters to “randomly se-

¹⁷The term “six degrees of separation” is not due to Milgram, but comes from a play written by John Guare in

lected” individuals in Omaha, Nebraska and Wichita, Kansas with instructions asking them to please pass the letter (and instructions) to a close friend of theirs who either knew or might be likely to know a particular doctor in Boston. Before doing so, they should also write their name on a roster to record the chain of message passing. Of the 64 letters that eventually reached the doctor—a small fraction of those sent out—the average length was only 5.5, not hundreds, and a legend was born.

Duncan Watts and Steve Strogatz, in a 1998 *Science* paper, studied this phenomenon using a toy model, now called the “small world model,” in which vertices are arranged on a 1-dimensional circular lattice (a “ring” network) and connected with their k nearest neighbors. Each edge is then rewired, with probability p , to connect a uniformly random pair of vertices. At $p = 0$, the network is fully ordered, where the density of local connections is largest and the diameter is $O(n)$. At $p = 1$, the network is fully disordered, local connections are absent and the diameter is $O(\log n)$. See Figure 2 below.




Figure 2: (A) Watts and Strogatz’s schematic illustrating the small worlds model. (B) The small-world phenomenon, in which at even modest levels of randomization, the network exhibits both small mean geodesic path lengths *and* a high clustering coefficient.

The interesting behavior emerges as we dial p between these two extremes: when only a small number of edges have been randomly rewired (small p), the diameter of the network collapses from $O(n)$ to $O(\log n)$ while the local structure is still largely preserved. That is, a highly-ordered “big world” can be transformed, by rewiring only a small number of connections, into a still mostly ordered small world, in which geodesic paths traverse a vanishing fraction of the network. This behavior reminded them of some properties of social networks, which have small diameter (as evidenced by

1990. The play was subsequently made into a movie of the same name starring Will Smith in 1993, and, ironically, not starring Kevin Bacon.

Milgram's study) but also many triangles.¹⁸

The small-world result is interesting for several reasons. First, it exemplifies an early effort to understand via a toy model how social networks can have both substantial local structure (here, many triangles) and a small diameter. Second, it shows that some measures of network structure can be extremely sensitive to uncertainty in the network structure. For instance, suppose we make a network using our toy model in which we have only rewired a few edges. If we, in fact, observe each of the network's edge with probability $p \ll 1$, then it is possible that we will miss the long-range connections that cross the ring. The result would be a substantial overestimate of the diameter of the network, potentially confusing $O(n)$ with $O(\log n)$. Fortunately, not all measures of network structure are this sensitive to sampling errors, but it is well worth our time to consider the impact of the data generation and data observation processes when analyzing and modeling networks.

2.5 Components

If every pair of vertices is connected by some path, then the network is *connected*. If there is some pair of vertices between which no path exists, the network is said to be *disconnected*, i.e., it is composed of more than one *component*.


In an undirected graph, the set of vertices reachable from one vertex is called a *component* (and, for every vertex j reachable from i , i is also reachable from j). The figure below shows an example of a set of components in an undirected graph.

In a directed graph, reachability in one direction does not imply reachability in the other, and the notion of “connected” becomes more nuanced. A group of nodes that is pairwise reachable only if we ignore the direction of the edges is a *weakly connected component*, while a group of nodes that is pairwise reachable if we obey the directions is a *strongly connected component*. Similarly, an *out component* is the set of vertices that can be reached from i , while an *in component* is the set of vertices that can reach i .

Many empirical networks are composed of multiple components, and among them, there is always a *largest connected component* (the LCC), which is usually the component of greatest interest. You can spot the LCC in the figure below quite easily. In mathematical models, the *giant component* is also the largest component, but we add an additional property, which is to say that the size of the giant component must be $O(n)$.¹⁹ We will revisit the notion of a giant component later in the

¹⁸The small worlds model was generalized to d -dimensions by Jon Kleinberg in 2000, who further showed that under a particular distribution of the long-range links, greedy routing is highly efficient and takes only $O(\log^2 n)$ steps in expectation. Some of my own early work showed that this result can be achieved dynamically and adaptively using a simple rewiring rule. There's a potential independent project here, if anyone is interested.

¹⁹The term “giant component” is meaningless in most empirical situations since it is only defined asymptotically.



semester, when we study models of random graphs.

Counting components.

To count and measure the size of the components within a network, we use any standard SSSP or APSP algorithm. For undirected networks, a breadth-first or depth-first search forest suffices, while for weighed or directed networks, Dijkstra’s algorithm works well.

While the algorithm runs, we need only label all vertices that are pairwise reachable with the same vertex label, in an auxiliary array. When the algorithm terminates, we may make a single pass through the array to count the number of unique labels (the number of components) and count the number of times each label occurs (the sizes of each component). In a weighted graph, identifying strongly connected components is most easily done via a depth-first search forest, in $O(n + m)$ time.

2.6 Other measures of network structure

These measures of network structure cover just the basics. Many other measures exist in the networks literature, and their creation often follows the needs of specific applications. You can think of each such measure as a function f that takes as input a graph G and outputs either a single

For empirical networks with multiple components, we instead focus on the largest component and reserve the term “giant component” for mathematical models.

scalar value or a vector of values. That is, each $f(G)$ projects the space of all possible graphs onto a low-dimensional space. As a result of this projection, information is surely lost and thus each measure $f(G)$ is an incomplete description of the original graph G . A good function f projects graphs in a way that highlights the structural feature of interest by creating good separation among points in the low-dimensional space.

By thinking about network measures as functions, it becomes reasonable to want a set of such functions $\{f\}$ that behave like an orthonormal basis set for the space of all graphs, or perhaps all graphs from a particular domain, say, all social networks. That is, are there a set of sufficient statistics to describe all possible structural patterns in a network? Unfortunately, at least at this point in the development of network science, the answer appears to be no, there is no such set.

Instead, what we do know is that many of the canonical network measures (many of which are described above, and summarized in Table 1) are highly correlated with each other. For instance, having a high degree tends to be inversely correlated with having a high local clustering coefficient, but having many high-degree vertices in a network tends to correlate with having a small diameter. Etc. This happens for a simple reason. Networks are not linearly decomposable. Instead, they are complicated, high-dimensional objects, and adding or removing even a single edge can, in some cases, have a dramatic impact on some network measures, but not others. For instance, the average degree is very robust to the addition of a few edges, but the diameter can be exquisitely sensitive (recall the Watts and Strogatz results of Section 2.4).

network measure	scope	graph	definition	explanation
degree	L	U	$k_i = \sum_{j=1}^n A_{ij}$	number of edges attached to vertex i
in-degree	L	D	$k_i^{\text{in}} = \sum_{j=1}^n A_{ji}$	number of arcs ending at vertex i
out-degree	L	D	$k_i^{\text{out}} = \sum_{j=1}^n A_{ij}$	number of arcs starting from vertex i
edge count	G	U	$m = \frac{1}{2} \sum_{ij} A_{ij}$	number of edges in the network
arc count	G	D	$m = \sum_{ij} A_{ij}$	number of arcs in the network
mean degree	G	U	$\langle k \rangle = 2m/n = \frac{1}{n} \sum_{i=1}^n k_i$	average number of edges per vertex
mean in- or out-degree	G	D	$\langle k^{\text{in}} \rangle = \langle k^{\text{out}} \rangle = 2m/n$	average number of in- or out-edges per vertex
reciprocity	G	D	$r = \frac{1}{m} \sum_{ij} A_{ij} A_{ji}$	fraction of arcs that are reciprocated
reciprocity	L	D	$r_i = \frac{1}{k_i} \sum_j A_{ij} A_{ji}$	fraction of arcs from i that are reciprocated
clustering coefficient	G	U	$c = \frac{\sum_{ijk} A_{ij} A_{jk} A_{ki}}{\sum_{ijk} A_{ij} A_{jk}}$	the network's triangle density
clustering coefficient	L	U	$c_i = \sum_{j < k} A_{ij} A_{jk} A_{ki} / \binom{k_i}{2}$	fraction of pairs of neighbors of i that are also connected
diameter	G	U	$d = \max_{ij} \ell_{ij}$	length of longest geodesic path in an undirected network
mean geodesic distance	G	UD	$\ell = \frac{1}{Z} \sum_{ij} \ell_{ij}$	average length of a geodesic path
eccentricity	G	UD	$\epsilon_i = \max_j \ell_{ij}$	length of longest geodesic path starting from i

Table 1: A few standard measures of network structure. The “scope” of a measure indicates whether it describes a vertex-level (local, L) or network-level (global, G) pattern, and the “graph” column indicates whether it is calculated on an undirected (U) or directed (D) graph. In each definition, adjacency matrix notation is used and it is assumed that $A_{ij} = 1$ if the edge (i, j) exists in the network and $A_{ij} = 0$ otherwise. For the geodesic paths, we assume that the graph is connected.

3 The shape of (some) networks

With these tools in hand, we can now begin to analyze the shape of networks. The big table on the next page lists a number of different empirical networks, drawn from social, informational, technological, and biological systems, and their values along a few of these measures. It is worth noting that the field has progressed greatly since this table was published in 2003, but there are relatively few other or newer “big tables” showing a specific set of network statistics for a large and diverse set of real-world networks. As such, it remains an instructive list with many lessons to teach.

What should be immediately clear from this table is that network sizes vary enormously. Even with the social network category, the smallest is $n = 573$ (student relationships) while the largest is $n = 4.7 \times 10^7$ (telephone call graph), more than 80,000 times larger. And yet, despite this enormous difference in size, the mean degree of these two networks differs by less than a factor of 2. Similarly, the largest mean geodesic distance in a social network (16.01; student relationships) is nearly as large as the mean geodesic distance for the Altavista WWW network (16.18).

Figure 3 shows the within-class distributions of these simple measures (as box plots). These reveal that, indeed, social and information networks tend to be larger than technological networks, and that biological networks tend to be the smallest. The mean degrees are all similar, being generally less than 20, although social and information networks have the broadest distributions here. Interestingly, the mean geodesic distance (normalized by network size) shows the opposite pattern, with technological and biological networks showing the greatest variance. These patterns suggest

	network	type	n	m	z	ℓ
social	film actors	undirected	449 913	25 516 482	113.43	3.48
	company directors	undirected	7 673	55 392	14.44	4.60
	math coauthorship	undirected	253 339	496 489	3.92	7.57
	physics coauthorship	undirected	52 909	245 300	9.27	6.19
	biology coauthorship	undirected	1 520 251	11 803 064	15.53	4.92
	telephone call graph	undirected	47 000 000	80 000 000	3.16	
	email messages	directed	59 912	86 300	1.44	4.95
	email address books	directed	16 881	57 029	3.38	5.22
	student relationships	undirected	573	477	1.66	16.01
	sexual contacts	undirected	2 810			
information	WWW nd.edu	directed	269 504	1 497 135	5.55	11.27
	WWW Altavista	directed	203 549 046	2 130 000 000	10.46	16.18
	citation network	directed	783 339	6 716 198	8.57	
	Roget's Thesaurus	directed	1 022	5 103	4.99	4.87
	word co-occurrence	undirected	460 902	17 000 000	70.13	
technological	Internet	undirected	10 697	31 992	5.98	3.31
	power grid	undirected	4 941	6 594	2.67	18.99
	train routes	undirected	587	19 603	66.79	2.16
	software packages	directed	1 439	1 723	1.20	2.42
	software classes	directed	1 377	2 213	1.61	1.51
	electronic circuits	undirected	24 097	53 248	4.34	11.05
	peer-to-peer network	undirected	880	1 296	1.47	4.28
biological	metabolic network	undirected	765	3 686	9.64	2.56
	protein interactions	undirected	2 115	2 240	2.12	6.80
	marine food web	directed	135	598	4.43	2.05
	freshwater food web	directed	92	997	10.84	1.90
	neural network	directed	307	2 359	7.68	3.97

Figure 3: Basic statistics for a number of published networks. The properties are: type of graph, directed or undirected, total number of vertices n , total number of edges m , mean degree z and mean geodesic distance ℓ . Table adapted from M.E.J. Newman, “The structure and function of complex networks.” *SIAM Review* **45**, 167–256 (2003).




Figure 4: Boxplots showing the within-class distributions of network size ($\log n$), mean degree, and mean geodesic distance (normalized by network size).

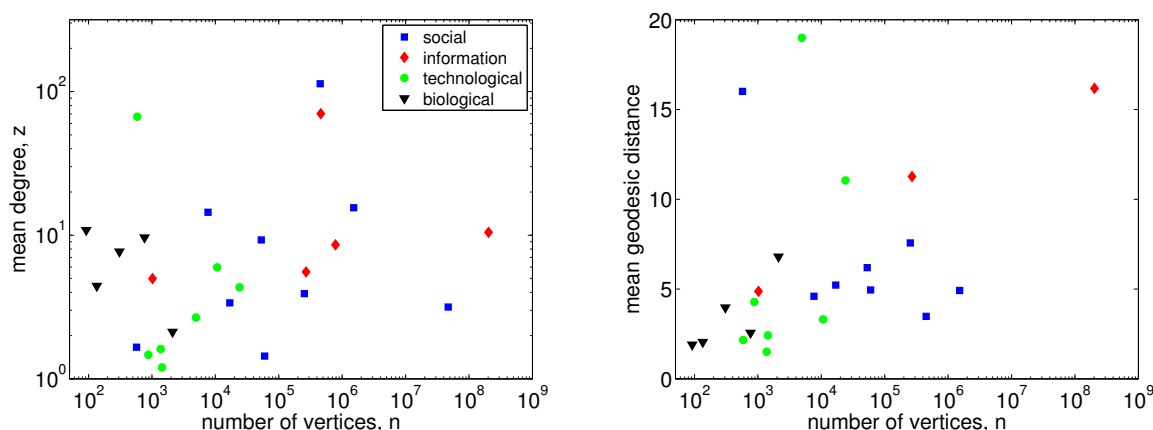


Figure 5: Scatter plots for mean degree and mean geodesic distance versus number of vertices.

that there are some broad patterns across these types of networks.

Another simple way to investigate the hypothesis that there is little overall variation in these measures across network class is to plot the mean degree and the mean geodesic distance versus network size. Figure 3 (next page) shows the results, with each class of network (social, information, technological, biological) given different colors and shapes. What is notable about these scatter plots is the lack of any clear pattern, either across the classes or within each class. That is, there appears to be little or no systematic correlation between mean degree and network size or between mean geodesic distance and network size.

4 At home

1. Peruse Chapters 1–5 (pages 15–104) in *Networks*
2. Read Chapter 6.1–6.12 (pages 109–149) in *Networks*